

# House Prices: Advanced Regression Techniques

Predict sales prices and practice feature engineering, RFs, and gradient boosting



December 2020

CSPB 3022 Final Project

Vidya Giri

Kaggle ID: vidyagi

Github Repository: <https://github.com/vidgi/predictingHousePrices>

## Table of Contents

1. Problem Description
2. Exploratory Data Analysis
3. Model Architecture
4. Results and Analysis
5. Final Conclusions

## 1. Problem Description

The Ames Housing dataset was compiled by Dean De Cock for use in data science education and was provided for use in this Kaggle prediction competition. There are 79 explanatory variables describing aspects of residential homes in Ames, Iowa.

The problem posed is to find an effective method to use the given explanatory variables to predict the final price of each home. To do so, we will need to understand the dataset given with the basic structure, conduct exploratory data analysis, build a predictive model that can be trained, and test our model's effectiveness and final results.

*Importing libraries and train and test data from csv files*

```
In [2]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import norm
from scipy import stats
%matplotlib inline
```

```
In [3]: train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

## Obtaining dimension/structure of data

```
In [4]: print('shapes:')
        print(train.shape)
        train.head()
```

```
shapes:
(1460, 81)
```

```
Out[4]:
```

|          | <b>Id</b> | <b>MSSubClass</b> | <b>MSZoning</b> | <b>LotFrontage</b> | <b>LotArea</b> | <b>Street</b> | <b>Alley</b> | <b>LotShape</b> | <b>LandContour</b> | <b>L</b> |
|----------|-----------|-------------------|-----------------|--------------------|----------------|---------------|--------------|-----------------|--------------------|----------|
| <b>0</b> | 1         | 60                | RL              | 65.0               | 8450           | Pave          | NaN          | Reg             |                    | Lvl      |
| <b>1</b> | 2         | 20                | RL              | 80.0               | 9600           | Pave          | NaN          | Reg             |                    | Lvl      |
| <b>2</b> | 3         | 60                | RL              | 68.0               | 11250          | Pave          | NaN          | IR1             |                    | Lvl      |
| <b>3</b> | 4         | 70                | RL              | 60.0               | 9550           | Pave          | NaN          | IR1             |                    | Lvl      |
| <b>4</b> | 5         | 60                | RL              | 84.0               | 14260          | Pave          | NaN          | IR1             |                    | Lvl      |

5 rows × 81 columns

## learning more about the data columns and their data types

```
In [5]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea              1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape             1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig             1460 non-null   object
11  LandSlope             1460 non-null   object
12  Neighborhood          1460 non-null   object
13  Condition1            1460 non-null   object
14  Condition2            1460 non-null   object
15  BldgType              1460 non-null   object
16  HouseStyle            1460 non-null   object
17  OverallQual           1460 non-null   int64
18  OverallCond           1460 non-null   int64
19  YearBuilt             1460 non-null   int64
20  YearRemodAdd          1460 non-null   int64
21  RoofStyle             1460 non-null   object
22  RoofMatl              1460 non-null   object
23  Exterior1st           1460 non-null   object
24  Exterior2nd           1460 non-null   object
25  MasVnrType            1452 non-null   object
26  MasVnrArea            1452 non-null   float64
27  ExterQual             1460 non-null   object
28  ExterCond             1460 non-null   object
29  Foundation            1460 non-null   object
30  BsmtQual              1423 non-null   object
31  BsmtCond              1423 non-null   object
32  BsmtExposure          1422 non-null   object
33  BsmtFinType1          1423 non-null   object
34  BsmtFinSF1            1460 non-null   int64
35  BsmtFinType2          1422 non-null   object
36  BsmtFinSF2            1460 non-null   int64
```

|    |               |      |          |         |
|----|---------------|------|----------|---------|
| 37 | BsmtUnfSF     | 1460 | non-null | int64   |
| 38 | TotalBsmtSF   | 1460 | non-null | int64   |
| 39 | Heating       | 1460 | non-null | object  |
| 40 | HeatingQC     | 1460 | non-null | object  |
| 41 | CentralAir    | 1460 | non-null | object  |
| 42 | Electrical    | 1459 | non-null | object  |
| 43 | 1stFlrSF      | 1460 | non-null | int64   |
| 44 | 2ndFlrSF      | 1460 | non-null | int64   |
| 45 | LowQualFinSF  | 1460 | non-null | int64   |
| 46 | GrLivArea     | 1460 | non-null | int64   |
| 47 | BsmtFullBath  | 1460 | non-null | int64   |
| 48 | BsmtHalfBath  | 1460 | non-null | int64   |
| 49 | FullBath      | 1460 | non-null | int64   |
| 50 | HalfBath      | 1460 | non-null | int64   |
| 51 | BedroomAbvGr  | 1460 | non-null | int64   |
| 52 | KitchenAbvGr  | 1460 | non-null | int64   |
| 53 | KitchenQual   | 1460 | non-null | object  |
| 54 | TotRmsAbvGrd  | 1460 | non-null | int64   |
| 55 | Functional    | 1460 | non-null | object  |
| 56 | Fireplaces    | 1460 | non-null | int64   |
| 57 | FireplaceQu   | 770  | non-null | object  |
| 58 | GarageType    | 1379 | non-null | object  |
| 59 | GarageYrBlt   | 1379 | non-null | float64 |
| 60 | GarageFinish  | 1379 | non-null | object  |
| 61 | GarageCars    | 1460 | non-null | int64   |
| 62 | GarageArea    | 1460 | non-null | int64   |
| 63 | GarageQual    | 1379 | non-null | object  |
| 64 | GarageCond    | 1379 | non-null | object  |
| 65 | PavedDrive    | 1460 | non-null | object  |
| 66 | WoodDeckSF    | 1460 | non-null | int64   |
| 67 | OpenPorchSF   | 1460 | non-null | int64   |
| 68 | EnclosedPorch | 1460 | non-null | int64   |
| 69 | 3SsnPorch     | 1460 | non-null | int64   |
| 70 | ScreenPorch   | 1460 | non-null | int64   |
| 71 | PoolArea      | 1460 | non-null | int64   |
| 72 | PoolQC        | 7    | non-null | object  |
| 73 | Fence         | 281  | non-null | object  |
| 74 | MiscFeature   | 54   | non-null | object  |
| 75 | MiscVal       | 1460 | non-null | int64   |
| 76 | MoSold        | 1460 | non-null | int64   |
| 77 | YrSold        | 1460 | non-null | int64   |
| 78 | SaleType      | 1460 | non-null | object  |
| 79 | SaleCondition | 1460 | non-null | object  |
| 80 | SalePrice     | 1460 | non-null | int64   |

dtypes: float64(3), int64(35), object(43)  
memory usage: 924.0+ KB

Overall, from looking at these properties of the data set we can see that there are 1460 rows of train data. From inspecting the data, we can see that there might be some columns that may have missing/not applicable data that we may need to filter out in the cleaning process.

Additionally, for this dataset there are 81 columns:

- House ID (ID)
- 36 are quantitative attributes (float64 and int64 data types)
- 43 are categorical attributes (object data types)
- Sale Price of the House (SalePrice)

Quantitative attributes: MSSubClass, LotFrontage, LotArea, OverallQual, OverallCond, YearBuilt, YearRemodAdd, MasVnrArea, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath, BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, TotRmsAbvGrd, Fireplaces, GarageYrBlt, GarageCars, GarageArea, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, MiscVal, MoSold, YrSold

Qualitative attributes: MSZoning, Street, Alley, LotShape, LandContour, Utilities, LotConfig, LandSlope, Neighborhood, Condition1, Condition2, BldgType, HouseStyle, RoofStyle, RoofMatl, Exterior1st, Exterior2nd, MasVnrType, ExterQual, ExterCond, Foundation, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinType2, Heating, HeatingQC, CentralAir, Electrical, KitchenQual, Functional, FireplaceQu, GarageType, GarageFinish, GarageQual, GarageCond, PavedDrive, PoolQC, Fence, MiscFeature, SaleType, SaleCondition


## 2. Exploratory Data Analysis

For exploratory data analysis, I will show the cleaning process, visualizations, histograms, and derive my plan of analysis.

Starting off, I wanted to plot the overall distribution of sales prices to see the general range of prices in our dataset. This does not appear to be a normal distribution and we will need to take into account our various features and attributes to understand their effect on the prices

```
In [6]: sns.histplot(train['SalePrice'],kde=True);  
        train['SalePrice'].describe()
```

```
Out[6]: count      1460.000000  
       mean      180921.195890  
       std       79442.502883  
       min       34900.000000  
       25%      129975.000000  
       50%      163000.000000  
       75%      214000.000000  
       max       755000.000000  
       Name: SalePrice, dtype: float64
```





To visualize the quantitative variables that might be significant further, I pairplotted the variables against SalePrice


```
In [7]: sns.pairplot(train,dropna=True, x_vars=['SalePrice'])
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x7ff805526520>
```









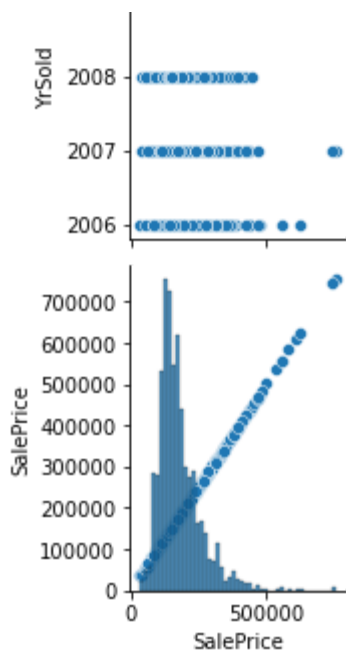













I also drew a heatmap to see any other variables that might show a significant relationship.

```
In [8]: corrmatrix = train.corr()
sns.heatmap(corrmatrix, cmap = 'coolwarm', vmax=.8, square=True)
```

Out[8]: <AxesSubplot:>



We can also check the correlation of the variables

```
In [9]: x = train.corr()
x = x[x['SalePrice']>0.5]
x['SalePrice']
```

```
Out[9]: OverallQual    0.790982
YearBuilt    0.522897
YearRemodAdd 0.507101
TotalBsmtSF  0.613581
1stFlrSF     0.605852
GrLivArea    0.708624
FullBath     0.560664
TotRmsAbvGrd 0.533723
GarageCars   0.640409
GarageArea   0.623431
```

```
SalePrice      1.000000
Name: SalePrice, dtype: float64
```

Some of the variables above appear to be correlated with SalesPrice as seen above where the plots generally seem to be positively correlated. The variables also appear to colinear since the variables are highly correlated, so we can use multi-linear regression. This can also be seen from the correlation table above.

Overall from this EDA we can start to narrow down some key variables that can come in handy for our model for now:

- OverallQual
- YearBuilt
- YearRemodAdd
- TotalBsmtSF
- GrLivArea
- FullBath
- TotRmsAbvGrd
- GarageCars
- GarageArea

### 3. Model Architecture

From the EDA, I can start to create a multilinear model using the variables I have derived above. First, I will construct a mult-linear model without interaction terms and predict the SalesPrice column on the other columns and print out the summary table.

Then, based on this, I will then construct a multi-linear model with interactions that are statistically significant at the  $p = 0.05$  level. I will start with full interactions and then eliminate interactions are do not meet the  $p = 0.05$  threshold to refine the model.

First, I will start off with all the variables with:

```
SalePrice~OverallQual*GrLivArea*GarageCars*TotalBsmtSF*FullBath*YearBuilt
```

This will enable me to check all of the interactions and narrow down the variables as needed.

```
In [20]: import statsmodels.formula.api as smf
import statsmodels.api as sm

model_trn = smf.ols(formula='SalePrice~OverallQual*GrLivArea*GarageCars*TotalBsmtSF*FullBath*YearBuilt', data=train_data)
# model_trn.summary() output omitted for report readability
```

This is a good start and I will now narrow down the variables even further by checking the individual probabilities to eliminate variables.

```
In [11]: model_trn = smf.ols(formula='SalePrice~OverallQual+GrLivArea+GarageCars+TotalBsmtSF+FullBath+YearBuilt', data=train_data)
model_trn.summary()
```

Out[11]:

#### OLS Regression Results

|                |                  |                     |         |
|----------------|------------------|---------------------|---------|
| Dep. Variable: | SalePrice        | R-squared:          | 0.772   |
| Model:         | OLS              | Adj. R-squared:     | 0.771   |
| Method:        | Least Squares    | F-statistic:        | 546.8   |
| Date:          | Tue, 08 Dec 2020 | Prob (F-statistic): | 0.00    |
| Time:          | 00:40:52         | Log-Likelihood:     | -17463. |

|                   |            |                   |           |       |           |           |
|-------------------|------------|-------------------|-----------|-------|-----------|-----------|
| No. Observations: | 1460       | AIC:              | 3.495e+04 |       |           |           |
| Df Residuals:     | 1450       | BIC:              | 3.500e+04 |       |           |           |
| Df Model:         | 9          |                   |           |       |           |           |
| Covariance Type:  | nonrobust  |                   |           |       |           |           |
|                   | coef       | std err           | t         | P> t  | [0.025    | 0.975]    |
| Intercept         | -1.159e+06 | 1.29e+05          | -8.978    | 0.000 | -1.41e+06 | -9.06e+05 |
| OverallQual       | 1.923e+04  | 1186.103          | 16.216    | 0.000 | 1.69e+04  | 2.16e+04  |
| GrLivArea         | 53.7230    | 4.158             | 12.920    | 0.000 | 45.567    | 61.879    |
| GarageCars        | 1.053e+04  | 3051.640          | 3.451     | 0.001 | 4543.945  | 1.65e+04  |
| TotalBsmtSF       | 28.9907    | 2.902             | 9.992     | 0.000 | 23.299    | 34.682    |
| FullBath          | -6408.3114 | 2685.771          | -2.386    | 0.017 | -1.17e+04 | -1139.899 |
| YearBuilt         | 258.0525   | 50.347            | 5.125     | 0.000 | 159.292   | 356.813   |
| YearRemodAdd      | 294.6839   | 63.791            | 4.620     | 0.000 | 169.552   | 419.816   |
| GarageArea        | 16.9253    | 10.314            | 1.641     | 0.101 | -3.307    | 37.158    |
| TotRmsAbvGrd      | 41.7519    | 1121.858          | 0.037     | 0.970 | -2158.886 | 2242.390  |
| Omnibus:          | 485.052    | Durbin-Watson:    | 1.980     |       |           |           |
| Prob(Omnibus):    | 0.000      | Jarque-Bera (JB): | 60868.329 |       |           |           |
| Skew:             | -0.439     | Prob(JB):         | 0.00      |       |           |           |
| Kurtosis:         | 34.620     | Cond. No.         | 4.42e+05  |       |           |           |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.42e+05. This might indicate that there are strong multicollinearity or other numerical problems.

From this, I can see that I will modify to the formula of

`SalePrice~OverallQual*GrLivArea*GarageCars+TotalBsmtSF+FullBath+YearBuilt`  
since I want to reduce the number of interactions to less than 0.05

```
In [12]: model_trn = smf.ols(formula='SalePrice~OverallQual*GrLivArea*GarageCars+TotalBsmtSF+FullBath+YearBuilt',
model_trn.summary())
```

|                          |                  |                            |           |
|--------------------------|------------------|----------------------------|-----------|
| OLS Regression Results   |                  |                            |           |
| <b>Dep. Variable:</b>    | SalePrice        | <b>R-squared:</b>          | 0.818     |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.816     |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 590.9     |
| <b>Date:</b>             | Tue, 08 Dec 2020 | <b>Prob (F-statistic):</b> | 0.00      |
| <b>Time:</b>             | 00:40:52         | <b>Log-Likelihood:</b>     | -17301.   |
| <b>No. Observations:</b> | 1460             | <b>AIC:</b>                | 3.463e+04 |
| <b>Df Residuals:</b>     | 1448             | <b>BIC:</b>                | 3.469e+04 |
| <b>Df Model:</b>         | 11               |                            |           |

| <b>Covariance Type:</b> |   | nonrobust  |                          |        |           |           |          |
|-------------------------|---|------------|--------------------------|--------|-----------|-----------|----------|
|                         |   | coef       | std err                  | t      | P> t      | [0.025    | 0.975]   |
|                         | <b>Intercept</b>                        | -1.075e+06 | 1.2e+05                  | -8.979 | 0.000     | -1.31e+06 | -8.4e+05 |
|                         | <b>OverallQual</b>                      | 2.54e+04   | 4318.728                 | 5.881  | 0.000     | 1.69e+04  | 3.39e+04 |
|                         | <b>GrLivArea</b>                        | 134.0080   | 17.511                   | 7.653  | 0.000     | 99.659    | 168.357  |
|                         | <b>OverallQual:GrLivArea</b>            | -18.7141   | 2.895                    | -6.464 | 0.000     | -24.393   | -13.035  |
|                         | <b>GarageCars</b>                       | 1.61e+04   | 1.12e+04                 | 1.433  | 0.152     | -5932.181 | 3.81e+04 |
|                         | <b>OverallQual:GarageCars</b>           | -2848.4830 | 1917.586                 | -1.485 | 0.138     | -6610.028 | 913.062  |
|                         | <b>GrLivArea:GarageCars</b>             | -41.7223   | 7.810                    | -5.342 | 0.000     | -57.043   | -26.402  |
|                         | <b>OverallQual:GrLivArea:GarageCars</b> | 8.8774     | 1.173                    | 7.570  | 0.000     | 6.577     | 11.178   |
|                         | <b>TotalBsmtSF</b>                      | 23.9166    | 2.598                    | 9.206  | 0.000     | 18.821    | 29.013   |
|                         | <b>FullBath</b>                         | -2980.1027 | 2414.626                 | -1.234 | 0.217     | -7716.641 | 1756.436 |
|                         | <b>YearBuilt</b>                        | 233.1520   | 45.135                   | 5.166  | 0.000     | 144.614   | 321.690  |
|                         | <b>YearRemodAdd</b>                     | 277.7551   | 57.379                   | 4.841  | 0.000     | 165.200   | 390.310  |
| <b>Omnibus:</b>         |   | 886.909    | <b>Durbin-Watson:</b>    |        | 2.000     |           |          |
| <b>Prob(Omnibus):</b>   |   | 0.000      | <b>Jarque-Bera (JB):</b> |        | 90081.683 |           |          |
| <b>Skew:</b>            |   | -1.920     | <b>Prob(JB):</b>         |        | 0.00      |           |          |
| <b>Kurtosis:</b>        |   | 41.289     | <b>Cond. No.</b>         |        | 3.74e+06  |           |          |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.74e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Finally, I will reduce the variables to OverallQual, GrLivArea, and YearBuilt.

```
In [13]: model_trn = smf.ols(formula='SalePrice~OverallQual*GrLivArea*YearBuilt', data=
model_trn.summary())
```

```
Out[13]:
```

|                          |                  |                            |           |
|--------------------------|------------------|----------------------------|-----------|
| OLS Regression Results   |                  |                            |           |
| <b>Dep. Variable:</b>    | SalePrice        | <b>R-squared:</b>          | 0.768     |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.767     |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 688.0     |
| <b>Date:</b>             | Tue, 08 Dec 2020 | <b>Prob (F-statistic):</b> | 0.00      |
| <b>Time:</b>             | 00:40:53         | <b>Log-Likelihood:</b>     | -17476.   |
| <b>No. Observations:</b> | 1460             | <b>AIC:</b>                | 3.497e+04 |
| <b>Df Residuals:</b>     | 1452             | <b>BIC:</b>                | 3.501e+04 |
| <b>Df Model:</b>         | 7                |                            |           |
| <b>Covariance Type:</b>  |                  | nonrobust                  |           |

|  | coef | std err | t | P> t | [0.025 | 0.975] |
|--|------|---------|---|------|--------|--------|
|--|------|---------|---|------|--------|--------|

|  |            |                          |           |       |           |           |
|--|------------|--------------------------|-----------|-------|-----------|-----------|
| <b>Intercept</b>                       | 4.031e+06  | 7.71e+05                 | 5.231     | 0.000 | 2.52e+06  | 5.54e+06  |
| <b>OverallQual</b>                     | -8.337e+05 | 1.19e+05                 | -7.004    | 0.000 | -1.07e+06 | -6e+05    |
| <b>GrLivArea</b>                       | -1639.8361 | 471.256                  | -3.480    | 0.001 | -2564.252 | -715.421  |
| <b>OverallQual:GrLivArea</b>           | 283.7362   | 64.683                   | 4.387     | 0.000 | 156.855   | 410.618   |
| <b>YearBuilt</b>                       | -2029.3788 | 392.755                  | -5.167    | 0.000 | -2799.806 | -1258.951 |
| <b>OverallQual:YearBuilt</b>           | 427.4303   | 60.382                   | 7.079     | 0.000 | 308.985   | 545.875   |
| <b>GrLivArea:YearBuilt</b>             | 0.8282     | 0.240                    | 3.447     | 0.001 | 0.357     | 1.300     |
| <b>OverallQual:GrLivArea:YearBuilt</b> | -0.1386    | 0.033                    | -4.220    | 0.000 | -0.203    | -0.074    |
| <b>Omnibus:</b>                        | 609.126    | <b>Durbin-Watson:</b>    | 1.966     |       |           |           |
| <b>Prob(Omnibus):</b>                  | 0.000      | <b>Jarque-Bera (JB):</b> | 48838.306 |       |           |           |
| <b>Skew:</b>                           | -1.034     | <b>Prob(JB):</b>         | 0.00      |       |           |           |
| <b>Kurtosis:</b>                       | 31.258     | <b>Cond. No.</b>         | 1.71e+10  |       |           |           |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.71e+10. This might indicate that there are strong multicollinearity or other numerical problems.

## 4. Result and Analysis

To check my results, I have created an errorcheck function and plotted the residuals to spot any outliers that we might want to remove to get a more accurate fit.



```
In [14]: def errorcheck(model):
    predicted = model_trn.predict(train)
    actual = train['SalePrice']
    actual = np.log(actual)
    predicted = np.log(predicted)
    error = np.sqrt(np.sum(np.square(actual - predicted)) / len(actual))
    return error

errorcheck(model_trn)
```

Out[14]: 0.191912740182196

Right now, we have an error of 0.178 and this can be reduced by removing some strategic points.

```
In [15]: sm.qqplot(model_trn.resid);
sm.graphics.plot_leverage_resid2(model_trn, alpha=0.05);
```



From this graph, we can see that outliers 185, 523, and 1298 might have high leverage over our model and can be removed as outliers. Let's inspect the data to ensure.

```
In [16]: train.iloc[[185,523,1298]]
```

```
Out[16]:
```

|             | <b>Id</b> | <b>MSSubClass</b> | <b>MSZoning</b> | <b>LotFrontage</b> | <b>LotArea</b> | <b>Street</b> | <b>Alley</b> | <b>LotShape</b> | <b>LandCon</b> |
|-------------|-----------|-------------------|-----------------|--------------------|----------------|---------------|--------------|-----------------|----------------|
| <b>185</b>  | 186       | 75                | RM              | 90.0               | 22950          | Pave          | NaN          | IR2             |                |
| <b>523</b>  | 524       | 60                | RL              | 130.0              | 40094          | Pave          | NaN          | IR1             |                |
| <b>1298</b> | 1299      | 60                | RL              | 313.0              | 63887          | Pave          | NaN          | IR3             |                |

3 rows × 81 columns

I decided to remove the data points and check if this improved my errorcheck result

```
In [17]: train.drop(1298, inplace=True)
train.drop(523, inplace=True)
train.drop(186, inplace=True)
```

So, I fit the model to the updated data frame and checked the error.

```
In [18]: model_trn = smf.ols(formula='SalePrice~OverallQual*GrLivArea*YearBuilt', data=train)
model_trn.summary()
errorcheck(model_trn)
```

```
Out[18]: 0.17816112832468034
```



And it showed an improvement from 0.178 to 0.166! Now we can use this updated model to create predictions on test.csv. I added some simple error checking to ensure that results were being generated for each data point:

```
In [19]: ypred = model_trn.predict(test)
         print(ypred)

         df = pd.read_csv('test.csv')
         output = pd.DataFrame({'Id': df.Id,
                                'SalePrice': ypred})
         output.to_csv('submission.csv', index=False)

         df = pd.read_csv('submission.csv')
         df = df.dropna()
         if(len(df)!=len(ypred)):
             print('Error with model prediction, submission.csv generated with errors')
         else:
             print('submission.csv generated')

0          118635.166827
1          150251.984093
2          154839.672112
3          183756.773356
4          203848.387351
...
1454         115610.353089
1455         115610.353089
1456         129568.974067
1457         135841.693880
1458         241077.112494
Length: 1459, dtype: float64
submission.csv generated
```

This generated the csv that I was able to submit to the Kaggle competition and got a score of 0.18716. This was not as high as some of the more complex models on Kaggle but, it was interesting to read more about what other people did to solve this problem and analyze the given data.

## 5. Final Conclusions

Overall, it was very interesting working through this problem and using skills from this class to generate a model that can be used for practical purposes. I am sure well-known websites like Zillow also use similar datasets to provide estimates of houses as well as other items such as car insurance, other types of insurance, and other price estimators and it is interesting to think about how you can take more features into account to build a more accurate model. Although my model did not quite reach the top of the Kaggle leaderboard, I am interested to learn more in the future and use other forms of modelling, tools, and frameworks to build better estimators and more robust models.